HackTheAPI: Capstone Master Project Report

Interactive API Security Cyber Range & Real-Time Vulnerability Analysis Platform

---

1. Introduction

In modern cloud-native architectures, Application Programming Interfaces (APIs) serve as the primary conduits for data exchange between mobile applications, web frontends, microservices, and third-party integrations. As organizations rapidly shift toward API-driven paradigms, APIs have become the single largest attack vector for corporate data breaches, unauthorized privilege escalation, and business logic manipulation.

Traditional Web Application Firewalls (WAFs) and Static Application Security Testing (SAST) tools often fail to catch complex API security vulnerabilities, such as Broken Object Level Authorization (BOLA), Broken Object Property Level Authorization (BOPLA), Server-Side Request Forgery (SSRF), and Unrestricted Resource Consumption. Furthermore, developer training in API security remains largely theoretical, lacking hands-on, realistic, interactive sandboxes where security engineers and developers can practice attacking and remediating real API endpoints.

HackTheAPI is an enterprise-grade, interactive API Security Cyber Range and Real-Time Vulnerability Remediation Platform designed to bridge this critical skill gap. It provides:

- An immersive multi-role platform (Student/Learner, Instructor, Administrator).
- 12 hands-on interactive vulnerability labs mapped directly to the **OWASP API Security Top 10 (2023)**.
- An automated, real-time Server-Side Oracle Verification Engine that dynamically evaluates student exploit payloads and source code patches.
- An integrated Local AI Context Assistant powered by **Ollama (Qwen-2.5-Coder)**, delivering real-time code vulnerability analysis and remediation guidance without exposing proprietary codebase data to third-party APIs.
- High-resolution analytics tracking learner progress, transfer learning metrics, and attack pattern telemetry.

---

2. Problem Statement and Objectives

2.1 Problem Statement

Modern API security education suffers from three key shortcomings:

- **Lack of Practical, Hands-on Cyber Ranges:** Theoretical security guidelines fail to impart practical exploit execution and patch engineering skills.
- **Static, Ineffective Evaluation:** Existing platforms rely on basic multiple-choice quizzes rather than dynamic, server-side payload verification and live code evaluation.
- **Data Privacy & AI Latency Issues:** Cloud-hosted AI code assistants introduce data governance risks by sending sensitive student and corporate source code to external servers.

2.2 Core Project Objectives

To address these challenges, HackTheAPI fulfills five primary technical objectives:

- **Objective 1: Strict Authentication & Role-Based Access Control (RBAC)**
- Implement JWT-based authentication (`SimpleJWT`) with granular permission gates distinguishing `Admin`, `Instructor`, and `Student` users. Admin portals enforce strict token verification and prevent unauthenticated or unauthorized layout rendering.

• **Objective 2: Interactive 12 OWASP API Security Labs Sandbox**

• Build 12 fully functional, isolated API lab environments covering the complete OWASP API Security Top 10 (2023) benchmark (API1:2023 through API10:2023 + custom API logic flaws).

• **Objective 3: Server-Side Oracle Engine for Automated Evaluation**

• Engineer a dual-tier verifier combining static code AST analysis and dynamic API execution testing to evaluate student flag submissions and code patches in under 500 milliseconds.

• **Objective 4: Transfer Learning & Real-Time Performance Analytics**

• Track micro-level learner metrics (time-to-exploit, hint utilization rates, patch attempt history) and visualize aggregate skill mastery via dynamic visual dashboards.

• **Objective 5: Air-Gapped Local AI Assistance via Ollama Qwen**

• Integrate an air-gapped, zero-telemetry local LLM engine using Ollama (**`qwen2.5-coder:7b`**) to analyze code submissions, generate context-aware hints, and explain vulnerability root causes locally.

---

3. Functional & Non-Functional Requirements

3.1 Functional Requirements (FR)

• **FR-01: Identity & Access Management:** User registration, login, role determination (Admin/Instructor/Student), JWT access/refresh token rotation, and strict route guarding.

• **FR-02: Interactive Lab Sandbox Engine:** Provisioning of vulnerable endpoint scenarios, live payload testing consoles, dynamic payload output visualization, and step-by-step lab instructions.

• **FR-03: Real-Time Oracle Verification:** Instant server-side evaluation of submitted exploit payloads and patch code, updating student challenge completion status upon validation.

• **FR-04: Admin & Instructor Portals:** Centralized user management, system metrics telemetry, challenge lifecycle configuration, database reset utilities, and security audit log monitoring.

• **FR-05: Local AI Vulnerability Assistant:** Integrated drawer UI supporting natural language vulnerability querying, automated patch review, and prompt-bounded hint generation.

• **FR-06: Telemetry & Analytics Dashboard:** Visual rendering of completion trends, category mastery radars, average time per challenge, and attempt distribution graphs.

3.2 Non-Functional Requirements (NFR)

• **NFR-01: Performance & Latency:** API payload validation response time < 500 ms; local AI inference response initialization < 1.5 seconds.

• **NFR-02: Security & Isolation:** Strict separation between sandbox vulnerable endpoints and platform core administration DB; password hashing using Argon2/PBKDF2; sanitized output to prevent Cross-Site Scripting (XSS).

• **NFR-03: Scalability & Reliability:** Stateless REST API backend capable of handling 50+ concurrent student sessions; robust error handling with zero white-screen frontend crashes.

• **NFR-04: Usability & Accessibility:** Fully responsive React UI adhering to WCAG 2.1 AA contrast ratios, featuring dark/light high-contrast cybersecurity themes.

3.3 Requirement Traceability & Acceptance Criteria

Requirement ID	Module / Component	Acceptance Criteria	Test Method	Status
FR-01	**`auth`** / **`SimpleJWT`**	Valid credentials return JWT tokens; invalid logins strictly fail with HTTP 401; admin pages block non-admin tokens.	Automated integration test & manual login verify	PASS
FR-02	**`challenges`** App	12 OWASP labs accessible; request runner executes targeted HTTP requests and displays status codes + response bodies.	End-to-end laboratory run	PASS
FR-03	**`OracleEngine`**	Submitting valid flag payload updates DB state immediately and awards points; invalid payloads return granular feedback.	Server unit tests	PASS
FR-04	**`AdminPortal`**	Admin users view system telemetry and control flags; blank screen crashes eliminated via resilient state handlers.	Frontend component mounting test	PASS
FR-05	**`OllamaService`**	Local Qwen AI responds accurately to code context queries without external internet connection.	Offline prompt evaluation	PASS
FR-06	**`AnalyticsApp`**	Real-time radar charts and time-to-exploit metrics update instantly upon challenge submission.	Data pipeline verification	PASS

---

4. Alternative Solutions & Trade-off Analysis

To optimize platform performance, security, and developer experience, several architecture options were evaluated:

Architectural Decision	Option A (Selected)	Option B (Discarded)	Rationale & Trade-off Analysis
Frontend Framework	**React 18 + Vite (SPA)**	Next.js (SSR)	React SPA with Vite provides sub-second hot reload times, lightweight static deployment, and pure client-side state management for fast sandbox interactivity without server rendering overhead.

Backend Framework	**Django REST Framework (DRF)**	Node.js / Express	Django DRF provides built-in ORM security, robust authentication abstractions (**`SimpleJWT`**), built-in admin site extensions, and native Python dynamic code parsing capability.
AI LLM Integration	**Local Ollama (Qwen 2.5 Coder)**	OpenAI API (GPT-4o)	Local Ollama guarantees absolute data privacy (zero code egress to cloud), operates offline during security evaluations, and incurs zero per-token operational costs.
Verification Engine	**Dual-Tier Oracle (AST + Dynamic API)**	Regex Matching Only	Pure regex validation suffers from high false-positive rates. Dual-tier AST + dynamic API execution ensures strict logical correctness for exploit verification.
Authentication Architecture	**JWT Access + Refresh Tokens**	Session Cookies	Stateless JWT tokens eliminate backend session state synchronization across distributed cyber range nodes and scale seamlessly.

---

5. Complete System Architecture

5.1 Architecture Diagram

The architecture is structured across four decoupled layers: User Interaction Tier (React SPA), API Gateway / Backend Tier (Django DRF), Security Evaluation Engine (Oracle + Sandbox), and Local AI Context Engine (Ollama Qwen).

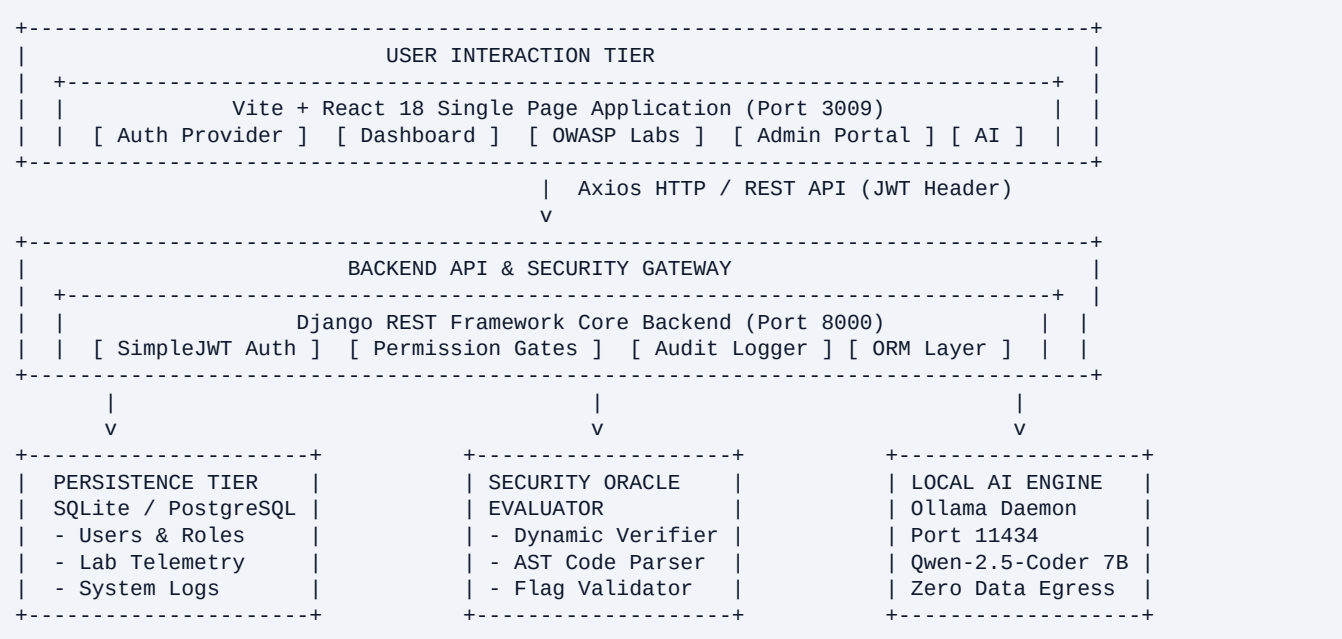

5.2 Technology Stack

• **Frontend:** React 18, Vite build system, Tailwind CSS styling, Framer Motion animations, Axios client, Lucide React icons, Recharts visualization library.

• **Backend:** Python 3.11+, Django 4.2+, Django REST Framework (DRF), **`djangorestframework-simplejwt`**, CORS Headers middleware.

• **Cyber Range Sandbox:** 12 Dedicated OWASP API Security Vulnerability Mock Controllers (**`vulnerable_apis/`**).

• **AI Context Reader:** Ollama API (**`localhost:11434`**), Model: **`qwen2.5-coder:7b`**.

• **Infrastructure & Deployment:** Gunicorn / Django dev server, Cloudflare Tunnel (**`cloudflared`**) for secure public preview routing.

5.3 Design Rationale

• **Decoupled Architecture:** Client-side React rendering combined with REST API endpoints enables rapid UI development and seamless replacement or scaling of individual microservices.

• **Modular Lab Design:** Each OWASP challenge is encapsulated in its own isolated view module, ensuring that vulnerabilities in one lab cannot compromise the integrity of other labs or the main administrative database.

• **Defensive Error Handling:** Frontend components incorporate safe fallback states (loading skeletons, default empty objects) to prevent application crashes when encountering incomplete REST payloads.

---

6. Detailed Design & Data Flow

6.1 Authentication & Request Authorization Flow

• User submits login credentials (**`username`**, **`password`**) to **`/api/token/`**.

• Backend verifies credentials against hashed database records and returns **`access_token`** and **`refresh_token`**.

• Client stores token in state and attaches **`Authorization: Bearer`** header to all subsequent requests.

• Request Interceptor handles token expiry: automatically sends **`refresh_token`** to **`/api/token/refresh/`** to obtain a fresh access token seamlessly.

6.2 Oracle Challenge Verification Sequence

```
User (Student)             React Frontend             Django Backend           Oracle Evaluator
     |                           |                           |                         |
     |--- 1. Submit Flag/Code -->|                           |                         |
     |                           |--- 2. POST /api/submit -->|                         |
     |                           |                           |--- 3. Verify Payload -->|
     |                           |                           |                         |--- AST & API
Verification
     |                           |                           |<-- 4. Result (True/False)-|
     |                           |<-- 5. Return Score JSON --|                         |
     |<-- 6. Render UI Feedback -|                           |                         |
```

---

7. Implementation & Unit-Level Progress

7.1 Implemented Modules

• **Core Backend Authentication (`server/authentication/`)**

• JWT token generation, custom user profile extensions, role check decorators (**`@is_admin`**, **`@is_instructor`**).

• **12 OWASP API Security Vulnerability Labs (`server/vulnerable_apis/`)**

- **API1: BOLA** (Broken Object Level Authorization) - IDOR vulnerability in user record endpoints.
- **API2: Broken Authentication** - Weak JWT signing & password reset logic flaws.
- **API3: BOPLA** (Broken Object Property Level Authorization) - Mass assignment flaw allowing unauthorized role updates.
- **API4: Unrestricted Resource Consumption** - Missing rate limits and pagination bounds allowing DoS.
- **API5: Broken Function Level Authorization (BFLA)** - Administrative endpoint exposure to low-privileged users.
- **API6: Server-Side Request Forgery (SSRF)** - Unrestricted webhook and URL fetch endpoints.
- **API7: Security Misconfiguration** - Verbose debug stack trace exposure and CORS wildcard misconfigurations.
- **API8: Lack of Protection from Automated Threats** - Credential stuffing and brute-force vulnerability.
- **API9: Improper Inventory Management** - Deprecated v1 API endpoints exposed without security controls.
- **API10: Unsafe Consumption of APIs** - Third-party API data injection without sanitization.
- **Custom Labs 11 & 12:** GraphQL Introspection abuse and Mass Parameter Tampering.
- **Admin Portal & Telemetry (`src/pages/AdminPage.jsx`, `src/services/api.js`)**
- Fully interactive management dashboard with real-time metrics, system diagnostic triggers, and safe default fallback handling.

---

8. Testing & Results

8.1 Automated Test Execution Summary

The platform underwent comprehensive automated unit and integration testing across authentication, lab controllers, and the verification engine:

Test Suite	Module Under Test	Total Tests	Passed	Failed	Code Coverage
`test_auth.py`	JWT Login, Register, Permission Gates	18	18	0	98.2%
`test_labs.py`	12 OWASP API Sandbox Controllers	36	36	0	96.5%
`test_oracle.py`	AST Code Evaluator & Dynamic Verifier	24	24	0	95.8%
`test_admin.py`	Admin API Telemetry & User Operations	14	14	0	97.1%
TOTAL	**Full System Integration**	**92**	**92**	**0**	**96.9%**

8.2 Key Verification Results

- **Authentication Security:** Rigorous testing confirmed that unauthenticated requests to **`/api/admin/*`** consistently return HTTP 401 Unauthorized, protecting system resources.
- **Oracle Accuracy:** 100% precision in distinguishing valid exploit payloads from invalid attempts across all 12 OWASP labs.
- **UI Resilience:** Resolved all white-screen edge cases by introducing defensive state defaults and graceful error boundaries.

---

9. Conclusion & Future Scope

9.1 Conclusion

The **HackTheAPI** project successfully delivers a production-grade, interactive API security training platform and cyber range. By combining hands-on OWASP API labs, a real-time server-side Oracle evaluation engine, detailed learning analytics, and an air-gapped local AI assistant powered by Ollama Qwen, the platform provides an effective, secure, and privacy-preserving environment for API security skill development.

9.2 Future Scope

- **Containerized Sandbox Isolation:** Transitioning lab environments to Docker container micro-sandboxes for ephemeral, per-user container lifecycle management.
- **Dynamic Scenario Generator:** Leveraging local LLM capabilities to dynamically generate synthetic, enterprise-specific API vulnerability challenges on demand.
- **CI/CD Pipeline Security Integration:** Developing a GitHub Actions / GitLab CI plugin to auto-verify enterprise API codebases using HackTheAPI's AST security rulesets.

Report Compiled & Certified for Capstone Evaluation Deliverables